

Programming Paradigms

Lecture 9

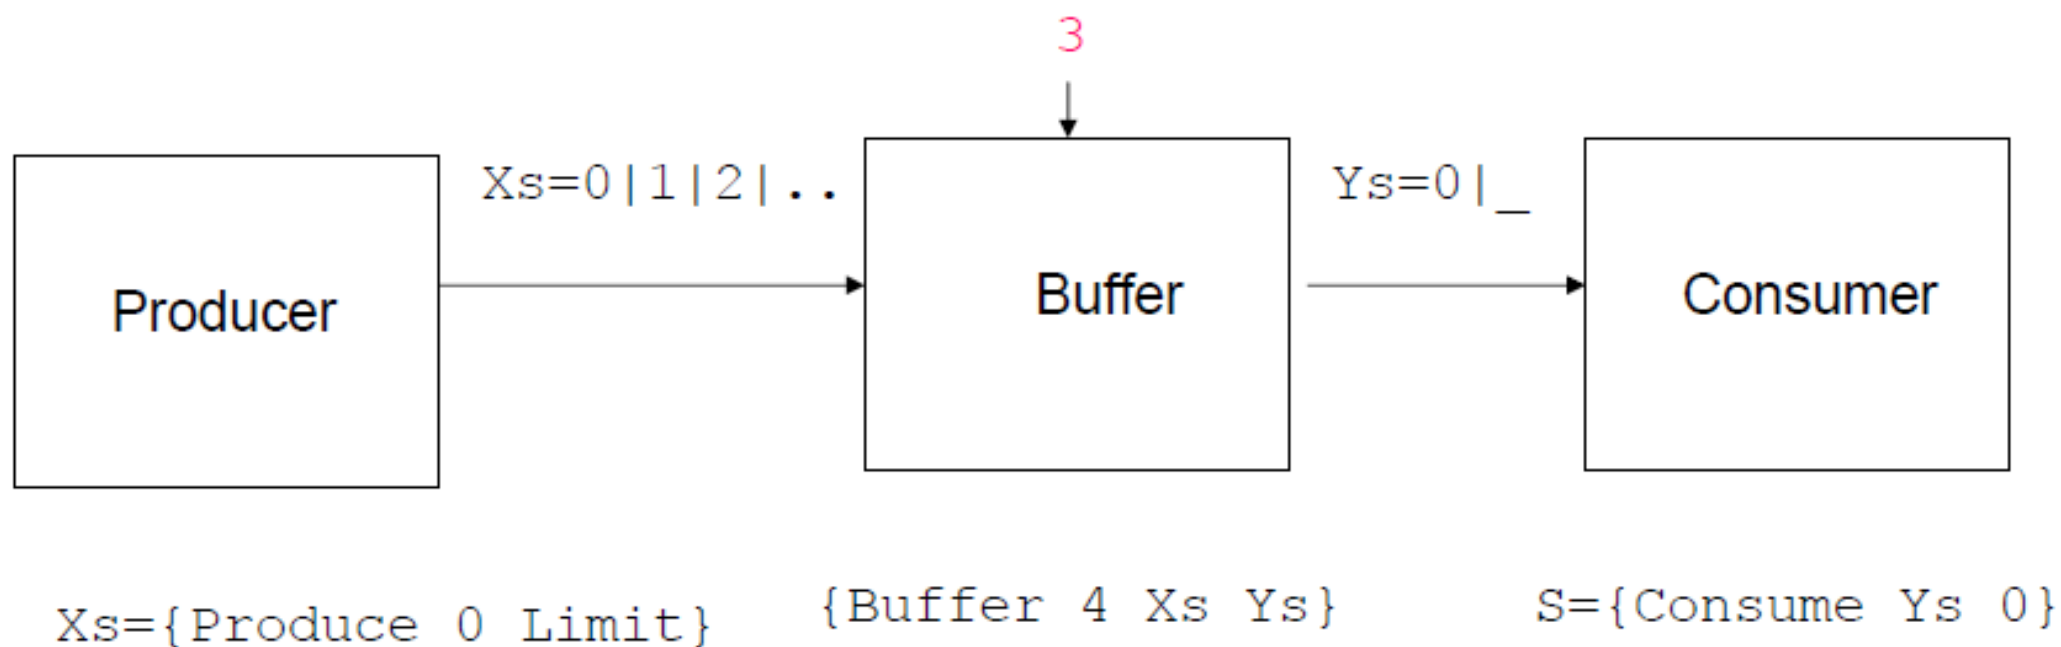
Slides are from Prof. Chin Wei-Ngan from NUS

More on Declarative Concurrency

Bounded Buffer

- *Eager – producer may run ahead*
 - *Demand-driven – consumer in control but more complex execution.*
 - *Compromise : Bounded Buffer*
-

Bounded Buffer



```
local Xs Ys S in
  thread {DGenerate 0 Xs} end
  thread {Buffer 4 Xs Ys} end
  thread S={DSum Ys 0 10} end
  {Browse Xs} {Browse Ys} {Browse S}
end
```

```
proc {DGenerate N Xs}  
  case Xs of X|Xr then  
    X=N  
    {DGenerate N+1 Xr}  
  end  
end
```

```
fun {DSum ?Xs A Limit}  
  if Limit > 0 then  
    X|Xr=Xs  
    in  
      {DSum Xr A+X Limit-1}  
    else A end  
end
```

```
proc {Buffer N ?Xs Ys}
  fun{Startup N ?Xs}
    if N==0 then Xs
    else Xr in Xs=_ | Xr {Startup N-1 Xr} end
  end
```

```
  proc {AskLoop Ys ?Xs ?End}
    case Ys of Y|Yr then Xr End2 in
      Xs=Y | Xr
      End=_ | End2
      {AskLoop Yr Xr End2}
    end
  end
```

```
  End={Startup N Xs}
in
  {AskLoop Ys Xs End}
end
```

Lazy Streams

- Better solution for demand-driven concurrency

Use Lazy Streams

That is consumer decides, so producer runs on request.

Needed Variables

- Idea:

- start execution,
- when value for variable needed
- suspend on the variable

- Value for variable needed...

...a thread suspends on variable!

Lazy Execution (Reminder)

- Up to now the execution order of each thread follows textual order.

Each statement is executed in order strict order, whether or not its results are needed later.

- This execution scheme is called *eager execution*, or *supply-driven* execution
- Another execution order is to execute each statement only if its results **are needed** somewhere in the program
- This scheme is called **lazy evaluation**, or **demand-driven evaluation**

Lazy Execution. Reminder

declare

fun lazy {F1 X} 2*X end

fun {F2 Y} Y*Y end

B = {F1 3}

{Browse B} → nothing (simply unbound B)

C = {F2 4}

{Browse C} → display 16

A = B+C → display 6 for B

- F1 is a lazy function
 - B = {F1 3} is executed only if its result is needed in A = B+C
-

Example

declare

fun lazy {F1 X} 2*X end

fun lazy {F2 Y} Y*Y end

B = {F1 3}

{Browse B} % \rightarrow nothing (simply unbound B)

C = {F2 4}

{Browse C} % \rightarrow nothing (simply unbound C)

- F1 and F2 are now lazy functions
- B = {F1 3} and C = {F2 4} are executed only if their results are needed in an expression, like: $A = B+C$


Example

```
declare
fun lazy {F1 X} 2*X end
fun lazy {F2 Y} Y*Y end
B = {F1 3}
{Browse B} % → display 6
C = {F2 4}
{Browse C} % → display 16
A = B+C
```

- F1 and F2 are now lazy functions
- B = {F1 3} and C = {F2 4} are executed because their results are needed in A = B+C


Example

- In **lazy execution**, an operation suspends until its result is needed
- Each suspended operation is triggered when another operation needs the value for its arguments
- In general, multiple suspended operations can start concurrently



Example II

- In **data-driven execution**, an operation suspends until the values of its arguments results are available
- In general, the suspended computation can start concurrently



Triggers

- A by-need trigger is a pair (F, X) :

- a zero-argument function F
- a variable X

- Trigger creation

$X = \{\text{ByNeed } F\}$ or equivalently

$\{\text{ByNeed } (\text{proc } \{\$ A\} A = \{F\} \text{ end}) X\}$

- If X is needed, then $X = \{\text{ByNeed } F\}$ means:

execute **thread** $X = \{F\}$ **end**

delete trigger, X becomes a normal variable

Example 1: ByNeed

`X = {ByNeed fun { $ } 4 end}`

- Executing `{Browse X}`
 - Shows: `X` (meaning not yet triggered)
 - `Browse` does not need the value of `X`
- Executing `T : Z = X + 1`
 - `X` is needed
 - current thread `T` blocks (`X` is not yet bound)
 - new thread created that binds `X` to `4`
 - thread `T` resumes and binds `Z` to `5`

Example 2: ByNeed

```
declare
fun {F1 X} {ByNeed fun {$} 2*X end} end
fun {F2 Y} {ByNeed fun {$} Y*Y end} end
B = {F1 3}
{Browse B} % simply display B
C = {F2 4}
{Browse C} % simply display C
```

Example 2: ByNeed

```
declare
fun {F1 X} {ByNeed fun {$} 2*X end} end
fun {F2 Y} {ByNeed fun {$} Y*Y end} end
B = {F1 3}
{Browse B} % display 6
C = {F2 4}
{Browse C} % display 16
A = B+C
```

Example 3: ByNeed

```
thread X={ByNeed fun { $ } 3 end} end  
thread Y={ByNeed fun { $ } 4 end} end  
thread Z=X+Y end
```

- Considering that each thread executes atomically, there are six possible executions.
- For lazy execution to be declarative, all of these executions must lead to equivalent stores.
- The addition will wait until the other two triggers are created, and these triggers will then be activated.

Lazy Functions

```
fun lazy {Produce N}  
    N|{Produce N+1}  
end
```

can be implemented with by-need triggers

```
fun {Produce N}  
    {ByNeed fun {$} N|{Produce N+1} end}  
end
```

Lazy Production

```
fun lazy {Produce N}  
    N | {Produce N+1}  
end
```

- Intuitive understanding: function executes only, if its output is needed

Example: Lazy Production

```
fun lazy {Produce N}  
    N|{Produce N+1}  
end  
declare Ns={Produce 0}  
{Browse Ns}
```

- Shows again Ns
 - Remember: `Browse` does not need the values of the variables
-

Example: Lazy Production

```
fun lazy {Produce N}  
    N|{Produce N+1}  
end  
declare Ns={Produce 0}
```

- Execute `_ = Ns.1`
 - needs the variable `Ns`
 - Browser now shows `0|_ or 0|<Future>`

Example: Lazy Production

```
fun lazy {Produce N}  
    N|{Produce N+1}  
end  
declare Ns={Produce 0}
```

- Execute `_ = Ns.2.2.1`
 - needs the variable `Ns.2.2`
 - Browser now shows `0|1|2|_`

Everything can be Lazy!

- Not only producers, but also transducers can be made lazy
 - Sketch
 - consumer needs variable
 - transducer is triggered, needs variable
 - producer is triggered
-

Lazy Transducer. Example

```
fun lazy {Inc Xs}  
  case Xs  
  of X|Xr then X+1 | {Inc Xr}  
  end  
end  
  
declare Xs={Inc {Inc {Produce N} } }
```

Global Summary

- Declarative concurrency
 - Mechanisms of concurrent program
 - Streams
 - Demand-driven execution
 - execute computation, if variable needed
 - need is suspension by a thread
 - requested computation is run in new thread
 - By-Need triggers
 - Lazy functions
-

More on Concurrency

Overview

- Stream Object
 - Thread Module and Composition
 - Soft Real-Time Programming
 - Agents and Message Passing
 - Protocols
 - Erlang
-

Stream Object

Diagram illustrating the Stream Object structure and its recursive definition:

```
proc {StreamObject S1 X1 ?T1}
  case S1 of M|S2 then N X2 T2 in
    {NextState M X1 N X2}
    T1 = N|T2 {StreamObject S2 X2 T2}
  [] nil then T1=nil end
end
```

$\text{StreamObject} :: [A], B, [C] \rightarrow ()$ $\text{NextState} :: A, B, C, A \rightarrow ()$
--

```
declare S0 X0 T0
```

```
thread {StreamObject S0 X0 T0} end
```

Thread Operations

Common Operations on Thread

<code>{Thread.this}</code>	return thread id
<code>{Thread.state T}</code>	return current state of T
<code>{Thread.suspend T}</code>	suspend T
<code>{Thread.resume T}</code>	resume T
<code>{Thread.preempt T}</code>	preempt T
<code>{Thread.terminate T}</code>	terminate T
<code>{Thread.injectException T}</code>	raise E in thread T
<code>{Thread.setPriority T P}</code>	set priority of T
<code>{Thread.setThisPriority P}</code>	set priority of thread

Common Property Operations


`{Property.get priorities}` get current priority ratios

`{Property.put priorities
 p (high:X medium:Y) }` set system priority ratios

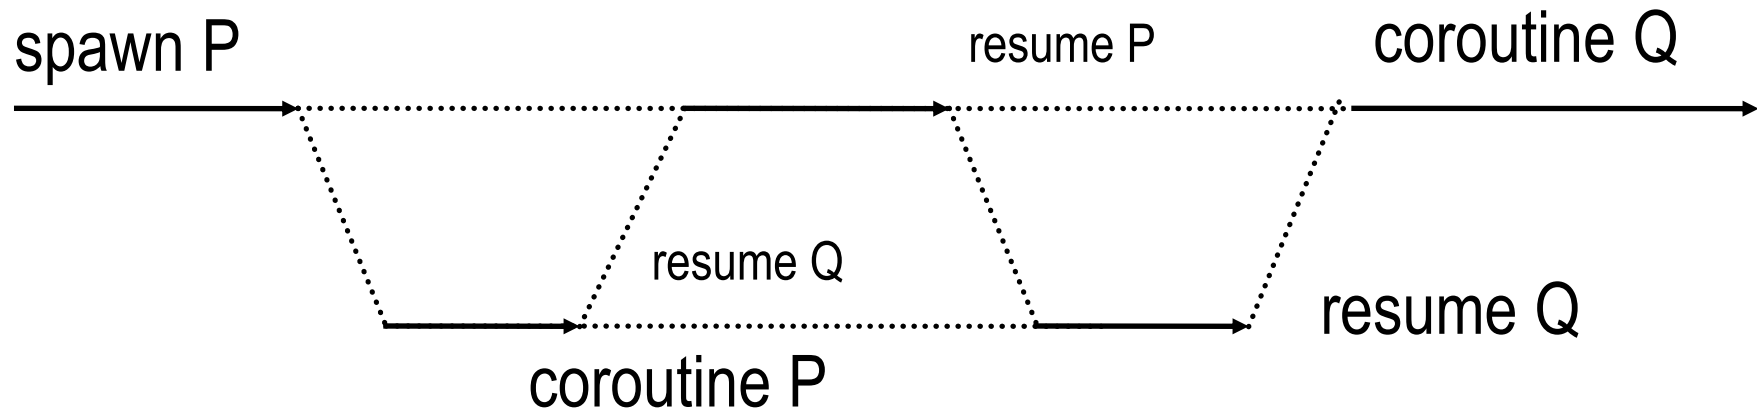
Coroutines

- Languages that do not support concurrent thread might instead support a notion called **coroutining**
- A coroutine is a nonpreemptive thread (sequence of instructions), there is no scheduler
- Switching between threads is the programmer's responsibility


Coroutines



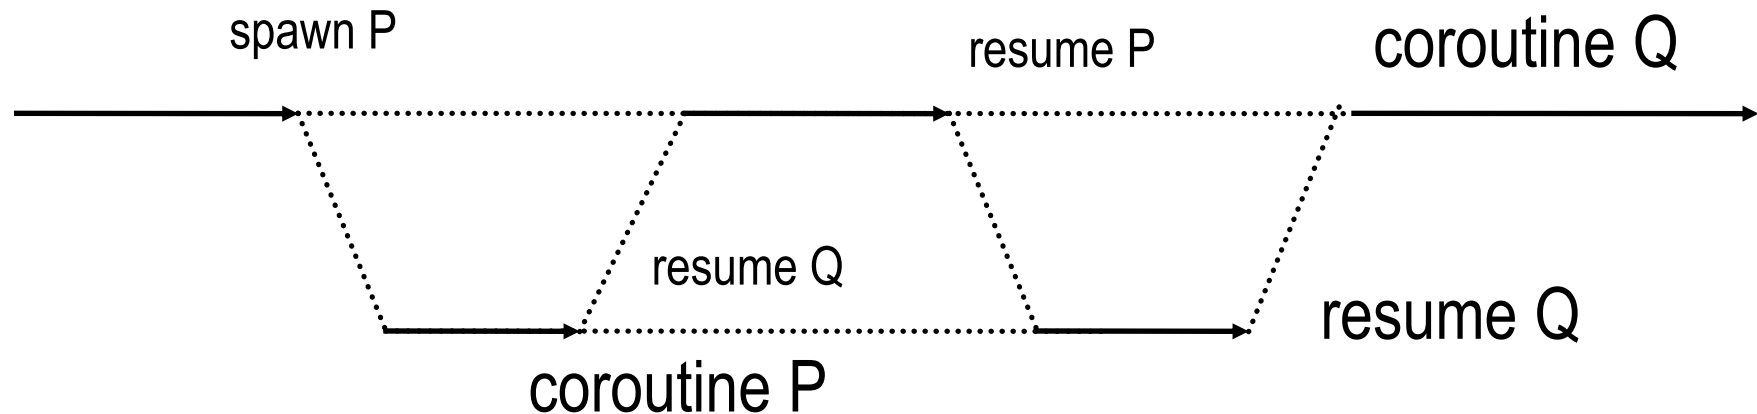
Procedures: one sequence of instructions, program transfers explicitly when terminated it returns to the caller



Coroutines

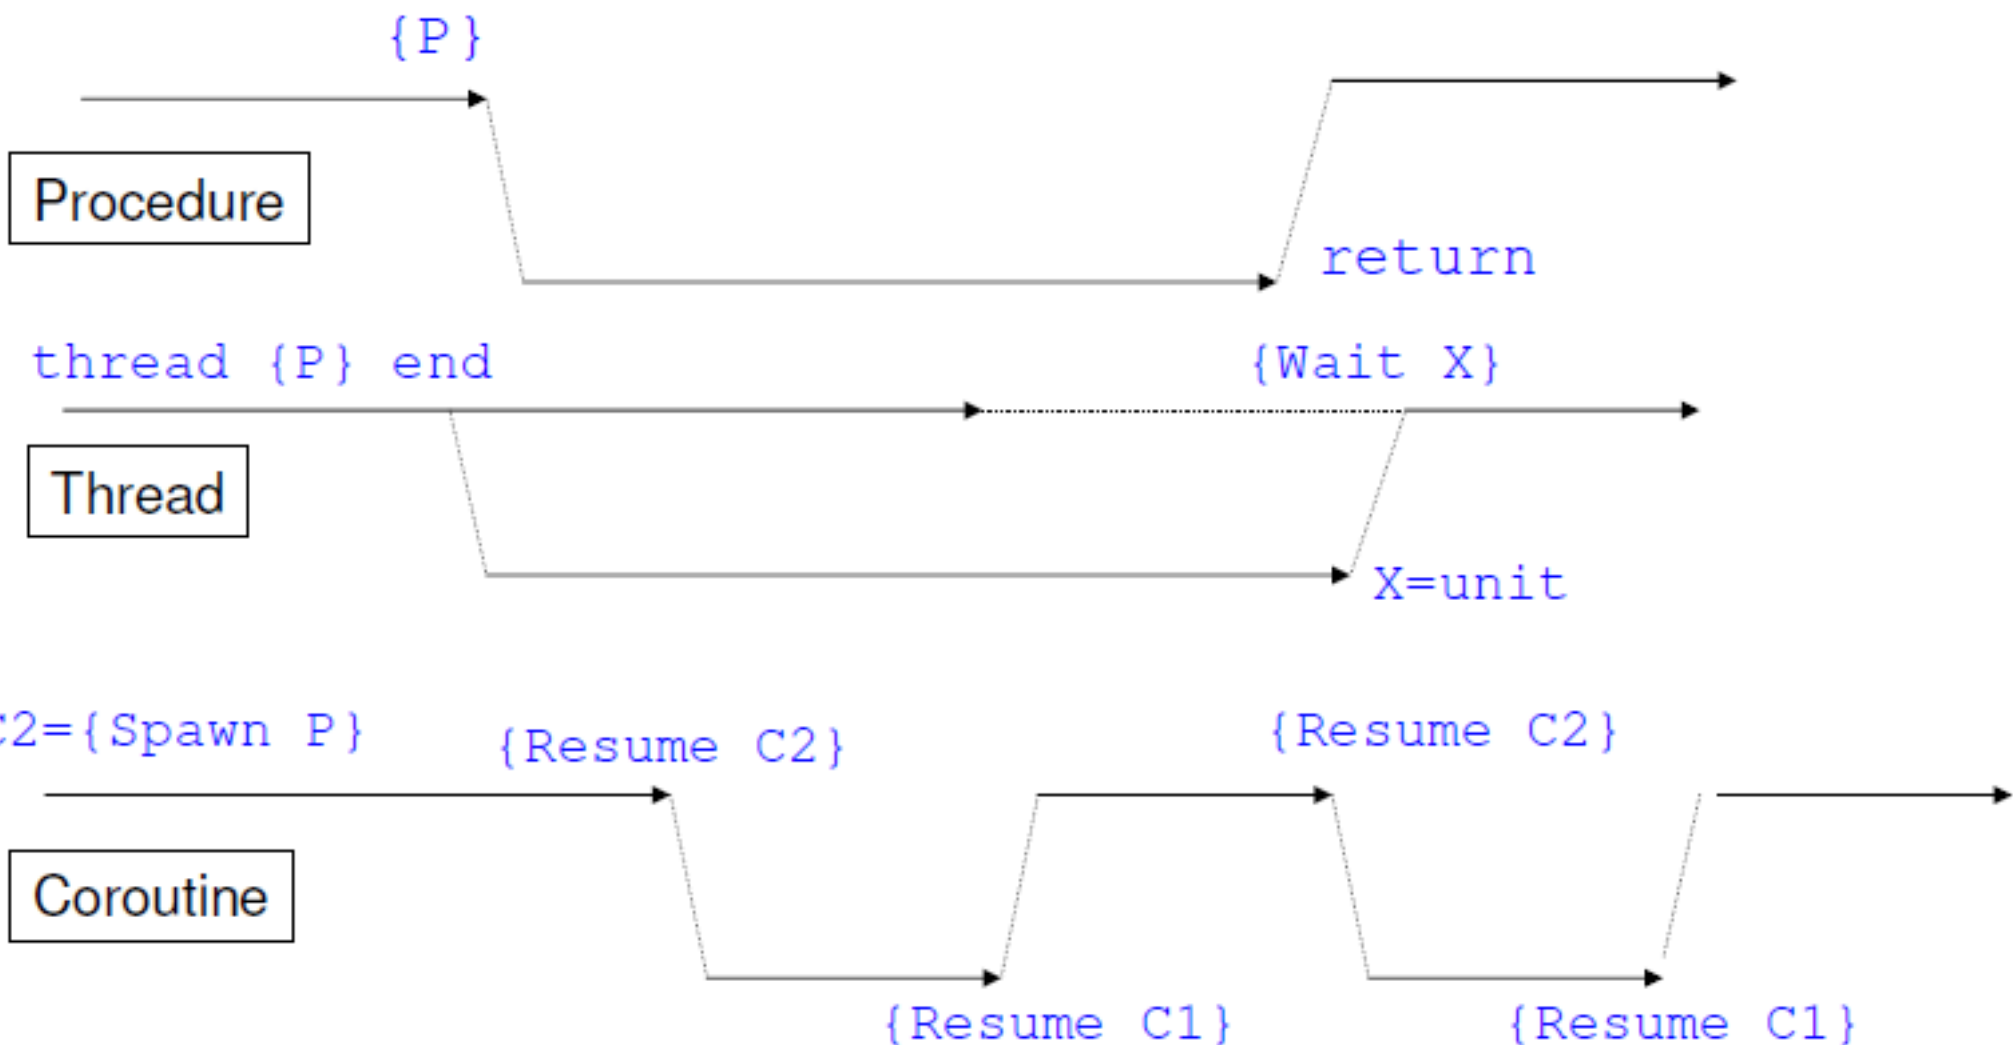


Coroutines: New sequences of instructions, programs explicitly does all the scheduling, by spawn, suspend and resume



Coroutine

A coroutine is a nonpreemptive thread



Basic Mechanism for Coroutines


```
fun {Spawn P}  
  PId in  
  thread  
    Pid={Thread.this}  
    {Thread.suspend Pid}  
    {P}  
  end  
  PId  
end
```

$\text{Spawn} :: (() \rightarrow ()) \rightarrow \text{Id}$ $\text{Resume} :: \text{Id} \rightarrow ()$
--

```
proc {Resume Id}  
  {Thread.resume Id}  
  {Thread.suspend {Thread.this}}  
end
```

Fork-Join for Threads

```
local  $X_1$   $X_2$  ..  $X_{n-1}$   $X_n$  in  
  thread <stmt1>  $X_1$ =unit end  
  thread <stmt2>  $X_2$ = $X_1$  end  
    :  
  thread <stmtn>  $X_n$ = $X_{n-1}$  end  
  {Wait  $X_n$ }  
end
```



wait for all threads to complete through variable binding

Barrier Synchronization

list of threads

```
proc {Barrier Ps}  
  fun {Loop Ps L}  
    case Ps of P|Pr then M in  
      thread {P} M=L end  
      {Loop Pr M}  
    [] nil then L  
  end  
end  
S={Loop Ps unit}  
in  
  {Wait S}  
end
```

wait for all threads to complete

Soft Real-Time Programming

- Real-time
 - control computations by time
 - animations, simulations, timeouts, ...
 - *Hard* real-time has firm deadlines, which have to be respected all the time, without any exception (medical equipments, air traffic control, ...)
 - *Soft* real-time is used in less demanding situations.
 - suggested time
 - no time guarantees
 - no hard deadlines as for controllers, etc.
 - Examples: telephony, consumer electronics, ...
-

The `Time` module

- The `Time` module contains a number of useful soft real-time operations:
 - `Delay`
 - `Alarm`
 - `Time`
 - `{Delay N}` suspends the thread for `N` milliseconds
 - Useful for building abstractions
 - timeouts
 - repeating actions
-

The `Time` module

- `{Alarm N U}` creates a new thread that binds `U` to `unit` after at least `N` milliseconds.
 - `Alarm` can be implemented with `Delay`
 - `{Time.time}` returns the integer number of seconds that have passed since the current year started
-

Soft Real-Time Programming. Example

```
functor
import
  Browser(browse:Browse)
define
  proc {Ping N}
    if N == 0 then {Browse 'ping terminated'}
    else {Delay 500} {Browse ping} {Ping N - 1} end
  end
  proc {Pong N}
    {For 1 N 1
      proc {$ I} {Delay 600} {Browse pong} end }
    {Browse 'pong terminated'}
  end
in
  {Browse 'game started'}
  thread {Ping 6} end
  thread {Pong 6} end
end
```

Soft Real-Time Programming. Example



Agents and Message Passing Concurrency

Client-Server Architectures

- Server provides some service
 - receives message
 - replies to message
 - examples: web server, mail server, ...
 - Clients know address of server and use service by sending messages
 - Server and clients run independently
-

Client-Server Applications ...

- With declarative programming, it is impossible to write a client/server program where the server does not know which client will send the next message.
 - **Observable *nondeterministic* behavior**: the server can receive information in any order from two independent clients.
 - The server has **only** an input stream from which it reads commands.
-

The Message-Passing Concurrent Model

- Extends the declarative concurrent model by adding one new concept, an **asynchronous communication channel**.
 - Any client can send messages to the channel at any time and the server can read all the messages from the channel (no limitations).
 - A client/server program may give different results on different executions because the order of clients' sends is not fixed.
 - **Message-passing model** is *nondeterministic* and therefore no longer declarative.
-

Peer-to-Peer Architectures

- Similar to Client-Server:
 - every client is also a server
 - communicate by sending messages to each other
 - We call all these guys (client, server, peer)
agent
 - In [van Roy, Haridi; 2004] book, this is called
portObject
-

Common Features

■ Agents

- | | |
|---------------------|-----------------------------|
| □ have identity | mail address |
| □ receive messages | mailbox |
| □ process messages | ordered mailbox |
| □ reply to messages | pre-addressed return letter |

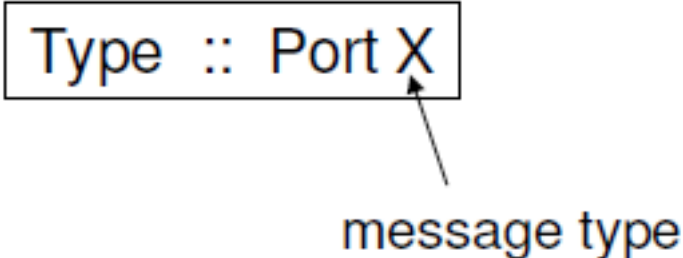
■ Now how to cast into programming language?

Message Sending

- Message data structure
- Address **port**
- Mailbox stream of messages
- Reply dataflow variable in message

Type :: Port X

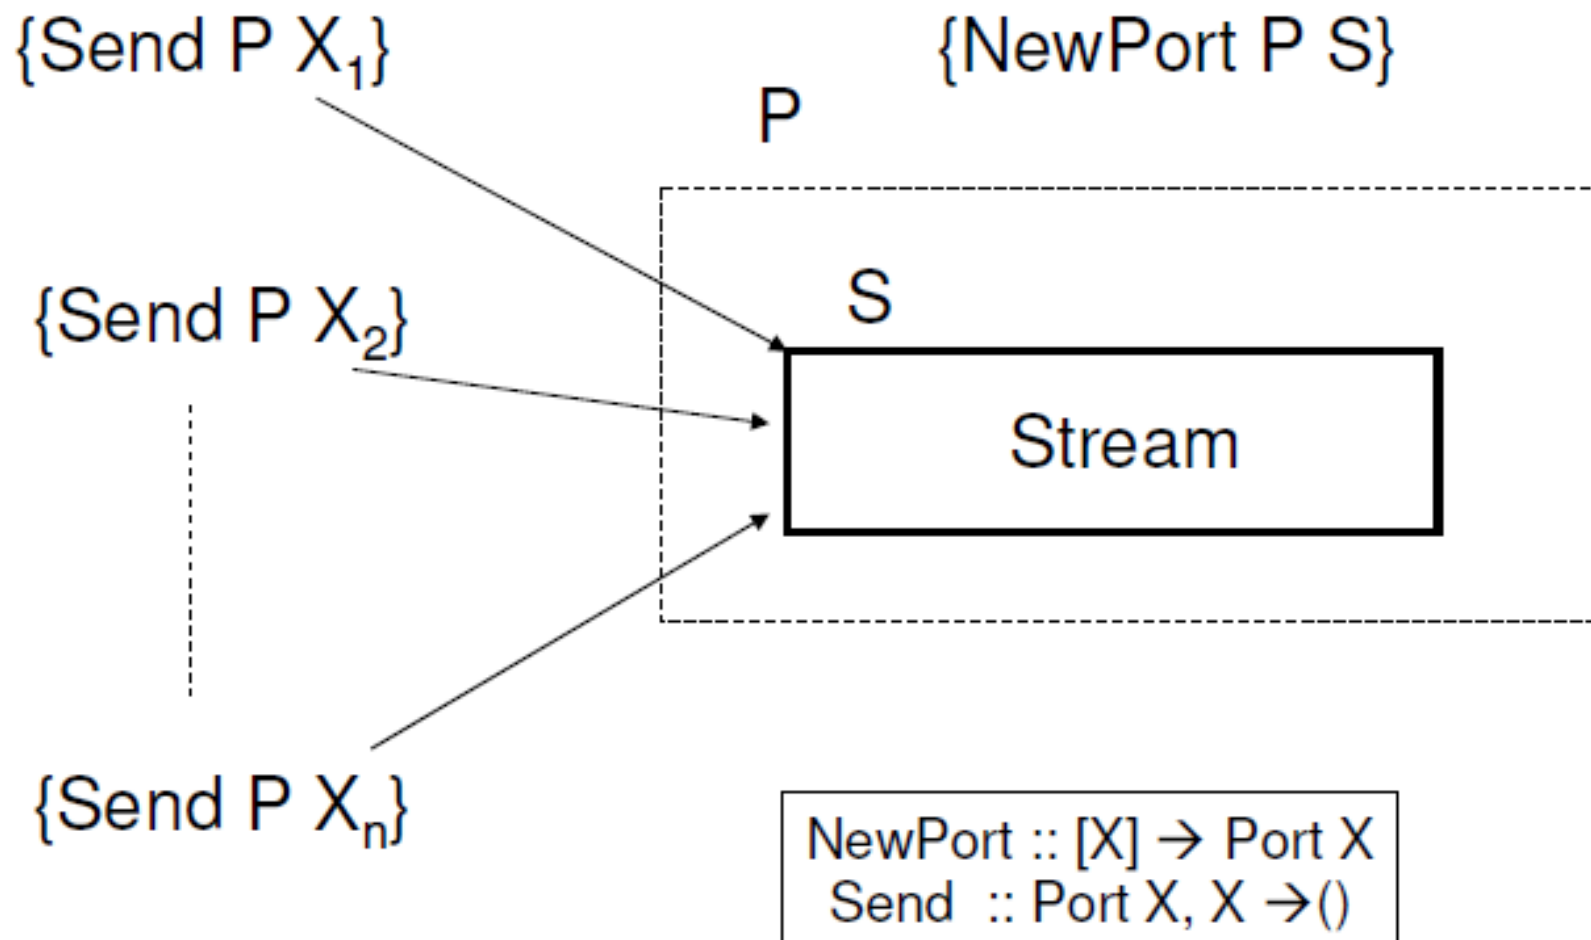
message type



Ports

- A **port** is an ADT with two operations:
 - $\{\text{NewPort } S \ P\}$ or equivalently $P = \{\text{NewPort } S\}$: create a new port with entry point (channel) P and stream S .
 - $\{\text{Send } P \ X\}$: append X to the stream corresponding to the entry point P .
 - Successive sends from the same thread appear on the stream in the same order in which they were executed.
 - This property implies that a port is an asynchronous FIFO (first-in, first-out) communication channel.
-

Port and its Stream



Ports

- **Asynchronous:** a thread can send a message at any time and it does not need to wait for any reply.
 - As soon as the message is in the communication channel, the thread can continue executing.
 - Communication channel can contain many pending messages, which are waiting to be handled.
-

Example

```
declare S P
```

```
P={NewPort S}
```

```
{Browse S}
```

- Displays initially `S<future>` (or `_`)

Example

```
declare S P
```

```
P={NewPort S}
```

```
{Browse S}
```

- **Execute** {Send P a}
- **Shows** a|_<future>

Example

```
declare S P
```

```
P={NewPort S}
```

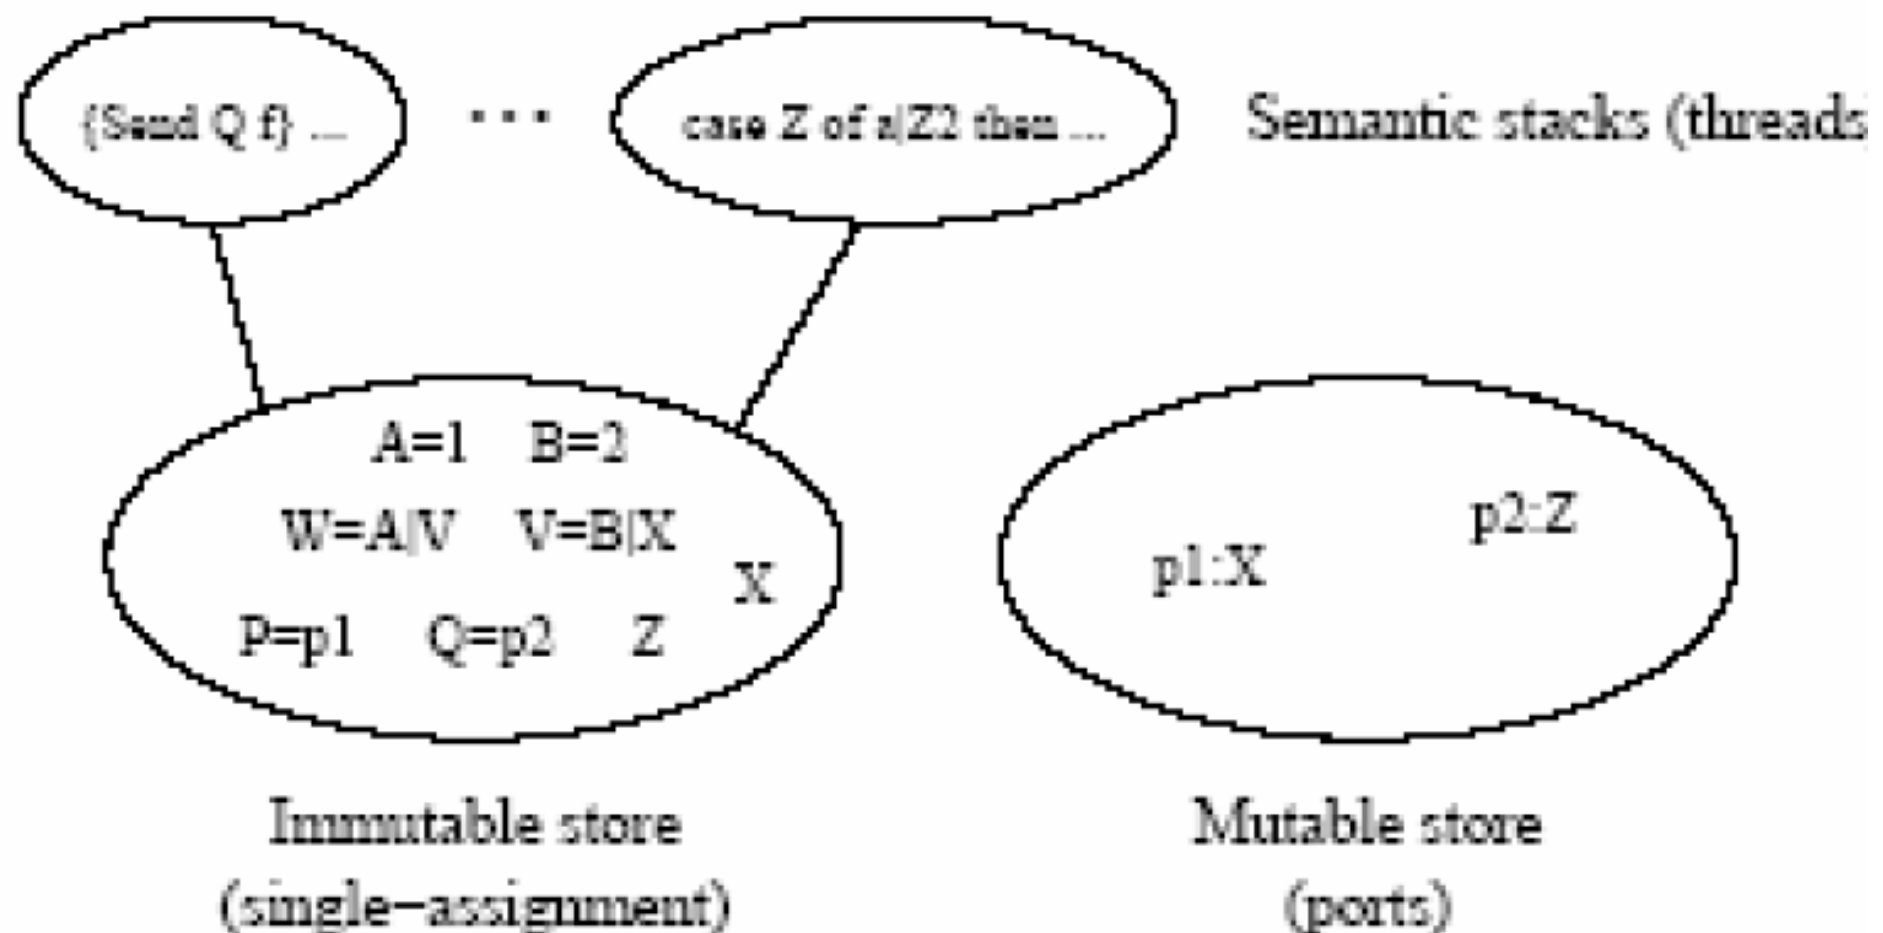
```
{Browse S}
```

- Execute `{Send P b}`
- Shows `a|b|_<future>`
- Note that `{Send P a}` and `{Send P b}` are in the same thread

Semantics of Ports

- Extend the execution state of the declarative model by adding a **mutable store** μ
- This store contains ports, i.e. pairs of the form $x : y$, where x and y are variables of the single-assignment store (x is the channel's name and y is the *current last position of stream*).
- The mutable store is initially empty.
- The semantics guarantees that x is *always bound* to a name value that represents a port and that y is *unbound*.
- The execution state becomes a triple (MST, σ, μ) (or (MST, σ, μ, τ) if the trigger store is considered).

The Message-Passing Concurrent Mode



The `NewPort` Operation

- The semantics of $(\{\text{NewPort } \langle x \rangle \langle y \rangle\}, E)$ is:
 - Create a **fresh port name** (also called **unique address**) n .
 - Bind $E(\langle y \rangle)$ and n in the store.
 - If the binding is successful, then add the pair $E(\langle y \rangle) : E(\langle x \rangle)$ to the mutable store μ .
 - If the binding fails, then raise an error condition.

The `send` Operation

- The semantics of $(\{\text{send } \langle x \rangle \langle y \rangle\}, E)$ is:
 - If the activation condition is `true` ($E(\langle x \rangle)$ is determined), then:
 - If $E(\langle x \rangle)$ is not bound to the name of a port, then raise an error condition.
 - If the mutable store contains $E(\langle x \rangle) : \mathbf{z}$, then:
 - Create a new variable **$\mathbf{z0}$** in the store.
 - Update the mutable store to be $E(\langle x \rangle) : \mathbf{z0}$.
 - Create a new list pair $E(\langle y \rangle) | \mathbf{z0}$ and bind \mathbf{z} with it in the store.
 - If the activation condition is `false`, then suspend execution.

Question

```
declare S P
P={NewPort S}
{Browse S}
thread {Send P a} end
thread {Send P b} end
```

- What will the Browser show?
- Note that each {Send P ...} is in a separate thread

Question

```
declare S P
P={NewPort S}
{Browse S}
thread {Send P a} end
thread {Send P b} end
```

- Which will the Browser show?
- Either
 - `a|b|_<future>` or
 - `b|a|_<future>`
- non-determinism: we can't say what

Answering Messages

- Traditional view
 - Include the entry port P' of the sender in the message:
`{Send P pair(Message P') }`
 - Receiver sends answer message to P'
`{Send P' AnsMessage }`
-

Answering Messages

- Do not reply by address, use something like pre-addressed reply envelope
 - dataflow variable!!!
 - `{Send P pair(Message Answer) }`
 - Receiver can bind `Answer!`
-