
Programming Paradigms

Lecture 10

Slides are from Prof. Chin Wei-Ngan from NUS

More on Declarative Concurrency

Agents and Message Passing Concurrency

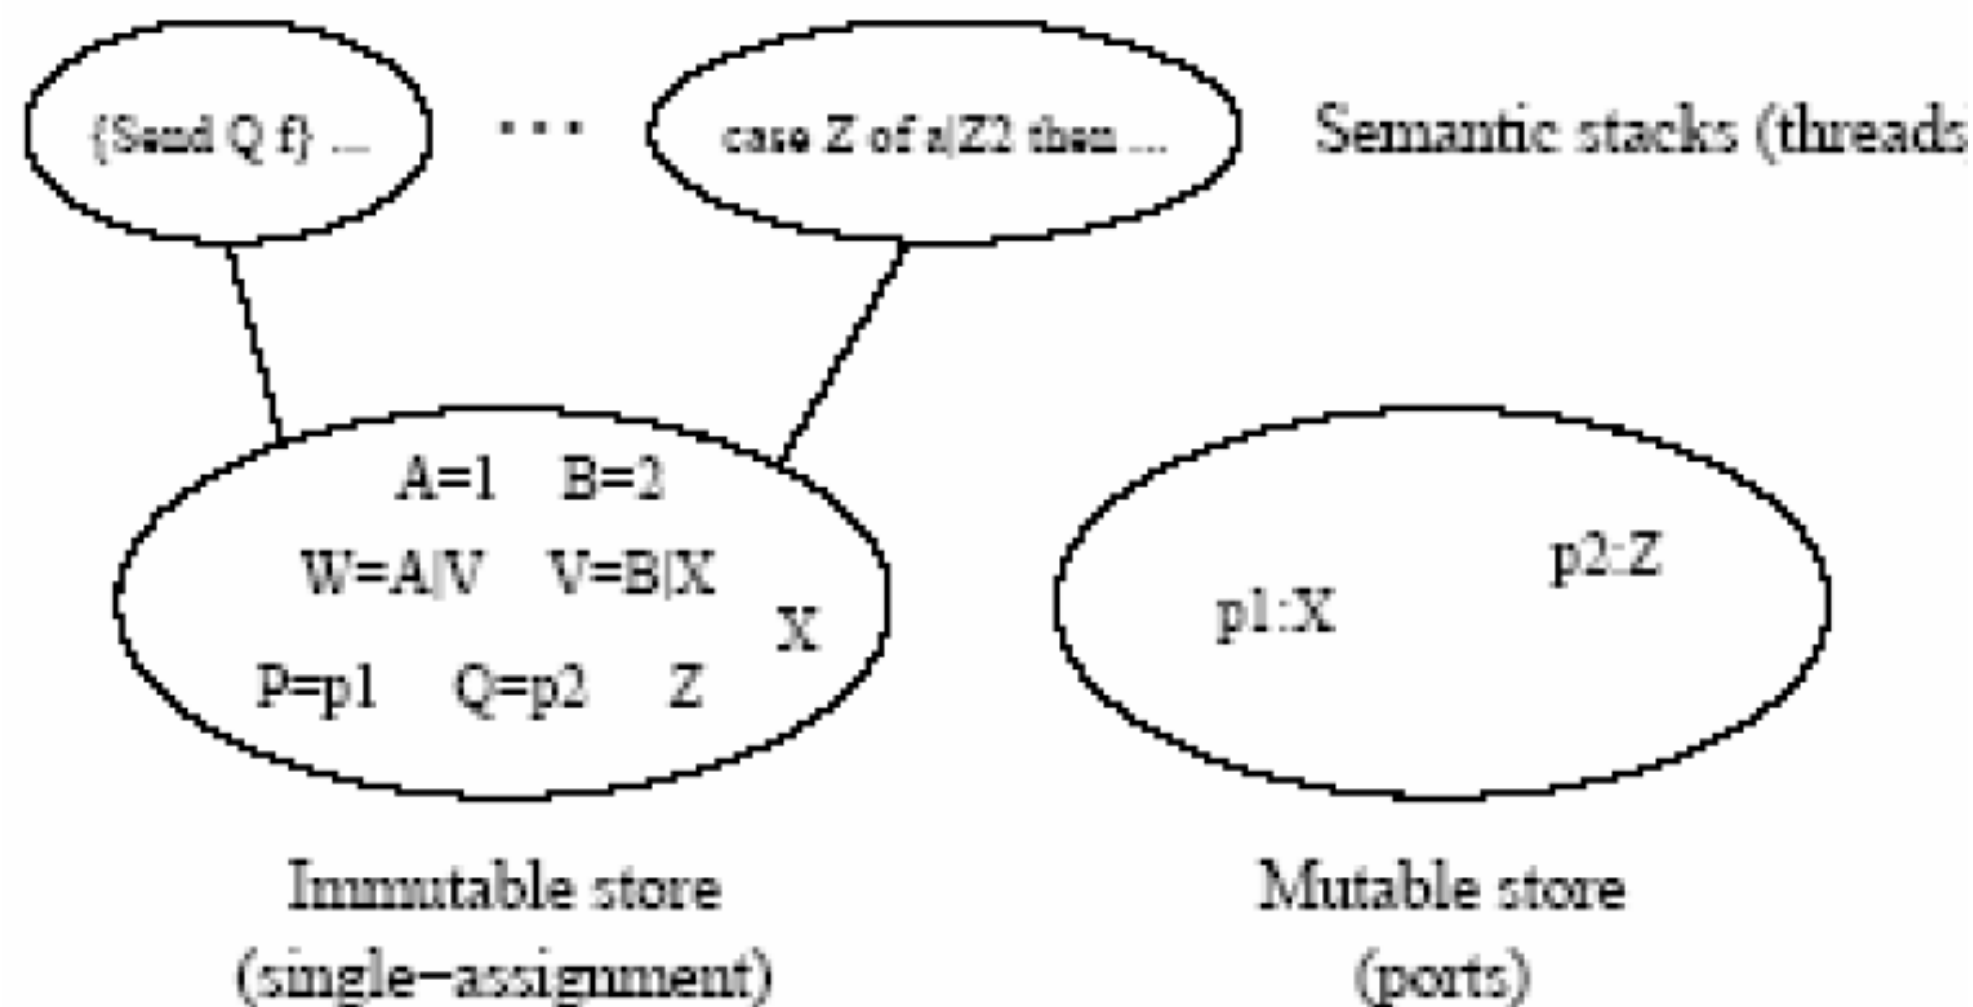
Ports

- A **port** is an ADT with two operations:
 - $\{\text{NewPort } S \ P\}$ or equivalently $P = \{\text{NewPort } S\}$: create a new port with entry point (channel) P and stream S .
 - $\{\text{Send } P \ X\}$: append X to the stream corresponding to the entry point P .
 - Successive sends from the same thread appear on the stream in the same order in which they were executed.
 - This property implies that a port is an asynchronous FIFO (first-in, first-out) communication channel.
-

Semantics of Ports

- Extend the execution state of the declarative model by adding a **mutable store** μ
- This store contains ports, i.e. pairs of the form $x : y$, where x and y are variables of the single-assignment store (x is the channel's name and y is the *current last position of stream*).
- The mutable store is initially empty.
- The semantics guarantees that x is *always bound* to a name value that represents a port and that y is *unbound*.
- The execution state becomes a triple (MST, σ, μ) (or (MST, σ, μ, τ) if the trigger store is considered).

The Message-Passing Concurrent Mode



The `NewPort` Operation

- The semantics of $(\{\text{NewPort } \langle x \rangle \langle y \rangle\}, E)$ is:
 - Create a **fresh port name** (also called **unique address**) n .
 - Bind $E(\langle y \rangle)$ and n in the store.
 - If the binding is successful, then add the pair $E(\langle y \rangle) : E(\langle x \rangle)$ to the mutable store μ .
 - If the binding fails, then raise an error condition.

The `send` Operation

- The semantics of $(\{\text{send } \langle x \rangle \langle y \rangle\}, E)$ is:
 - If the activation condition is `true` ($E(\langle x \rangle)$ is determined), then:
 - If $E(\langle x \rangle)$ is not bound to the name of a port, then raise an error condition.
 - If the mutable store contains $E(\langle x \rangle) : \mathbf{z}$, then:
 - Create a new variable **$\mathbf{z0}$** in the store.
 - Update the mutable store to be $E(\langle x \rangle) : \mathbf{z0}$.
 - Create a new list pair $E(\langle y \rangle) | \mathbf{z0}$ and bind \mathbf{z} with it in the store.
 - If the activation condition is `false`, then suspend execution.

Question

```
declare S P  
P={NewPort S}  
{Browse S}  
thread {Send P a} end  
thread {Send P b} end
```

- What will the Browser show?
- Note that each {Send P ...} is in a separate thread

Question

```
declare S P
P={NewPort S}
{Browse S}
thread {Send P a} end
thread {Send P b} end
```

- Which will the Browser show?
- Either
 - `a|b|_<future>` or
 - `b|a|_<future>`
- non-determinism: we can't say what

Answering Messages

- Traditional view
 - Include the entry port P' of the sender in the message:
`{Send P pair(Message P') }`
 - Receiver sends answer message to P'
`{Send P' AnsMessage }`
-

Answering Messages

- Do not reply by address, use something like pre-addressed reply envelope
 - dataflow variable!!!
 - `{Send P pair(Message Answer) }`
 - Receiver can bind `Answer!`
-

Port Objects

- A **port object** is a combination of one or more ports and a stream object.
 - This *extends stream objects* in two ways:
 - First, many-to-one communication is possible: many threads can reference a given port object and send to it independently.
 - This is not possible with a stream object because it has to know where its next message will come from.
 - Second, port objects can be embedded inside data structures (including messages).
 - This is not possible with a stream object because it is referenced by a stream that can be extended by just one thread.
-

Stream Object

Diagram illustrating the Stream Object structure and its recursive definition:

```
proc {StreamObject S1 X1 ?T1}
  case S1 of M|S2 then N X2 T2 in
    {NextState M X1 N X2}
    T1 = N|T2 {StreamObject S2 X2 T2}
  [] nil then T1=nil end
end
```

StreamObject :: [A], B, [C] → () NextState :: A,B, C,A → ()
--

```
declare S0 X0 T0
```

```
thread {StreamObject S0 X0 T0} end
```

Port Objects. Distributed Algorithm

```
declare P1 P2 ... Pn in  
local S1 S2 ... Sn in  
    {NewPort S1 P1}  
    {NewPort S2 P2}  
    ...  
    {NewPort Sn Pn}  
    thread {RP S1 S2 ... Sn} end  
end
```

- The thread contains a recursive procedure `RP` that reads the port streams and performs some action for each message received.
- Sending a message to the port object is just sending a message to one of the ports.
- Similar terms: **agent**, **process** (Erlang), **active object**

A Math Agent

```
proc {Math E}  
  case E  
  of add(N M Answer) then Answer=N+M  
  [] mul(N M Answer) then Answer=N*M  
  [] int(Formula Answer) then  
    Answer = ...  
  end  
end
```

- **Remark:** Answer is included in the stream's element X of {Send EntryPoint X}

Making the Agent Work (Port Creation)

```
local S in
  MP = {NewPort S}
  proc {MathProcess Ms}
    case Ms of M|Mr then
      {Math M}
      {MathProcess Mr}
    end
  end
  thread {MathProcess S} end
end
```

- `MathProcess` is a recursive procedure that reads the port streams and performs some action for each message received.

Making the Agent Work (Sending a Message)

```
declare A B
thread % client 1
    {Send MP add(2 3 A) }
    {Browse A}
end
thread % client 2
    {Send MP mul(2 3 B) }
    {Browse B}
end
```

- A and B are two dataflow variables which will be bound in port MP
-

Recall Higher-Order Construct

$\text{ForAll} :: \{[X], X \rightarrow ()\} \rightarrow ()$

```
proc {ForAll Xs P}  
  case Xs  
  of nil    then skip  
    [] X|Xr then {P X} {ForAll Xr P}  
  end  
end
```

- Call procedure P for all elements in X_S

Smells of Higher-Order...

- Using `ForAll`, we have

```
proc {MathProcess Ms}  
  {ForAll Ms Math}  
end
```

Making the Agent Work

```
declare MP in  
local S in  
    MP = {NewPort S}  
    thread {ForAll S Math} end  
end
```

Making the Agent Work

```
declare MP in  
local S in  
    MP = {NewPort S}  
    thread for M in S do {Math M} end end  
end
```

- The stream `s` is private (`local`) to the port.
 - `Math` is associated to the port `MP`
 - `MP` and `Math` can become arguments of a generic function.
-

Smells Even Stronger...

- Programming with port objects can be abstracted into a function

```
fun {NewAgent Process}
```

```
  Port Stream
```

NewAgent :: {X→()} → Port X

```
in
```

```
  Port={NewPort Stream}
```

```
  thread {ForAll Stream Process} end
```

```
  Port
```

```
end
```

- So, the previous port creation is equivalent with:

```
MP = {NewAgent Math}
```

Why Do Agents/ Processes Matter?

- Model to capture communicating entities
 - Each agent is simply defined in terms of how it replies to messages
 - Each agent has a thread of its own
 - no screw-up with concurrency
 - we can easily extend the model so that each agent has a state (encapsulated)
 - ***Extremely useful to model distributed systems!***
-

Summary so far

- Ports for message sending
 - use stream (list of messages) as mailbox
 - port serves as unique address
 - Use agent abstraction
 - combines port with thread running agent
 - simple concurrency scheme
 - Introduces non-determinism... and state!
-

Protocols

Protocols

- **Protocol:** is a set of rules for sending and receiving messages
 - programming with agents
 - Most well-known protocols:
 - the Internet protocols (TCP/IP, HTTP, FTP, etc.)
 - LAN (Local Area Network) protocols such as Ethernet and DHCP (Dynamic Host Connection Protocol), ...
-

RMI (Remote Method Invocation)

- It seems to be the most popular of the simple protocols.
 - It allows an object to call another object in a different operating system process, either on the same machine or on another machine connected by a network.
 - **RMI** is a descendant of the **RPC** (Remote Procedure Call), which was invented in 1980, before object-oriented programming became popular.
 - RMI became popular once objects started replacing procedures as the remote entities to be called.
 - We assume that a “*method*” is simply what a port object does when it receives a particular message.
-

Differences between RPC and RMI


RPC

- Faster than RMI
- Depends on the platform
- Has to convert the arguments between architectures so that each computer can use its native datatype

RMI

- Is part of Java's object-oriented approach
 - Allows multiple-concurrent method invocation
 - Is portable (doesn't depend on the operating system)
 - Good security system
 - To call outside methods, RMI needs JNI, JDBC, RMI-IIOP, RMI-IDL, etc.
-

1. RMI
(2 calls)



- A client sends a request to a server and then waits for the server to send back a reply.
- C stands for client, S for server, *idle* means “available to service requests”, *suspended* means “not available”.

The Server as a Port Object

```
declare
proc {ServerProc Msg}
  case Msg
  of calc(X Y) then
    Y = X * X + 1.0
  end
end
end
Server={NewAgent ServerProc}
```

- The second argument Y of `calc` is bound by the server.
- The server computes the polynomial $X * X + 1.0$

What is NewAgent? (Reminder)

```
fun {NewAgent Process}
  Port Stream
in
  Port={NewPort Stream}
  thread {ForAll Stream Process} end
  Port
end
```

The Client (using RMI)

```
declare
proc {ClientProc Msg}
  case Msg
  of work(Y) then
    Y1 Y2 in
      {Send Server calc(1.0 Y1)}
      {Wait Y1}
      {Send Server calc(2.0 Y2)}
      {Wait Y2}
      Y = Y1 + Y2
    end
  end
end
Client={NewAgent ClientProc}
```

The Client as a Port Object II

```
local X in
  {Send Client work(X) }
  {Browse X}
end
```

■ Difference between the client and server:

- ❑ The client definition references the server directly but the server definition does not know its clients.
 - ❑ The server gets a client reference indirectly, through the argument Y , i.e. the dataflow variable that is bound to the answer by the server.
 - ❑ The client waits until receiving the reply before continuing.
-

What is `Wait`?

- `{Wait X}` suspends the thread until `x` becomes determined, i.e. also called *explicit synchronization* on variable `x`

```
declare Y
```

```
{ByNeed proc {$ X} X=1 end Y}
```

```
{Browse Y}
```

```
{Wait Y}
```

```
<statement>
```

- Display `Y` in the Browser.
- To access `Y`, the operation `{Wait Y}` will trigger the producing procedure.
- `<statement>` will be executed only after `Y` is bound


Characteristics of RMI

- In the previous example, all messages are executed sequentially by the server.
 - In practice, some RMI implementations do things somewhat differently, i.e. they allow multiple calls from different clients to be processed concurrently.
 - May use different languages and different OS.
-

Asynchronous RMI

- Similar to RMI, except that the client continues execution immediately after sending the request.
 - The client is informed when the reply arrives.
 - So, two requests can be done in rapid succession.
 - **Motivation:** If communications between client and server are slow, then this will give a large performance advantage over RMI.
-

2. Asynchronous RMI (2 calls)



The Asynchronous RMI Client

```
declare
proc {ClientProc Msg}
  case Msg
  of work(Y) then Y1 Y2 in
    {Send Server calc(1.0 Y1)}
    {Send Server calc(2.0 Y2)}
    Y = Y1 + Y2
  end
end
Client={NewAgent ClientProc}
local X in
  {Send Client work(X)}
  {Browse X}
end
```

Characteristics of Asynchronous RMI

- Message sends overlap. Client waits for both results y_1 and y_2 before doing the addition $y_1 + y_2$.
- The server is the same as with standard RMI. It still receives messages one by one and executes them sequentially.
- Requests are handled by the server in the same order as they are sent and the replies arrive in that order as well.

RMI with Callback

- Server may need to call back client to fulfill request, e.g. check on some special values.

```
proc {ServerProc Msg}
  case Msg
  of calc(X ?Y Client) then X1 D in
    {Send Client delta(D)} ← callback
    X1=X+D
    Y = X * X + 1.0
  end
end
Server={NewAgent ServerProc}
```

RMI with Callback

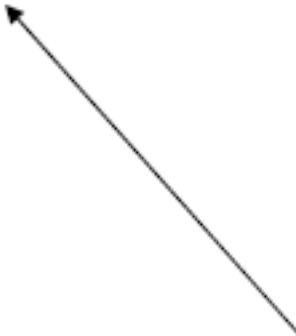
```
proc {ClientProc Msg}
  case Msg
  of work(?Z) then Y in
    {Send Server calc(10.0 Y Client)}
    Z=Y+100.0
  [] delta(?D) then D=1.0
  end
end
end
Client={NewAgent ClientProc}
{Browse {Send Client work($)}}}
```

■ Does this work?

No! It deadlocks as server and client waiting for each other.

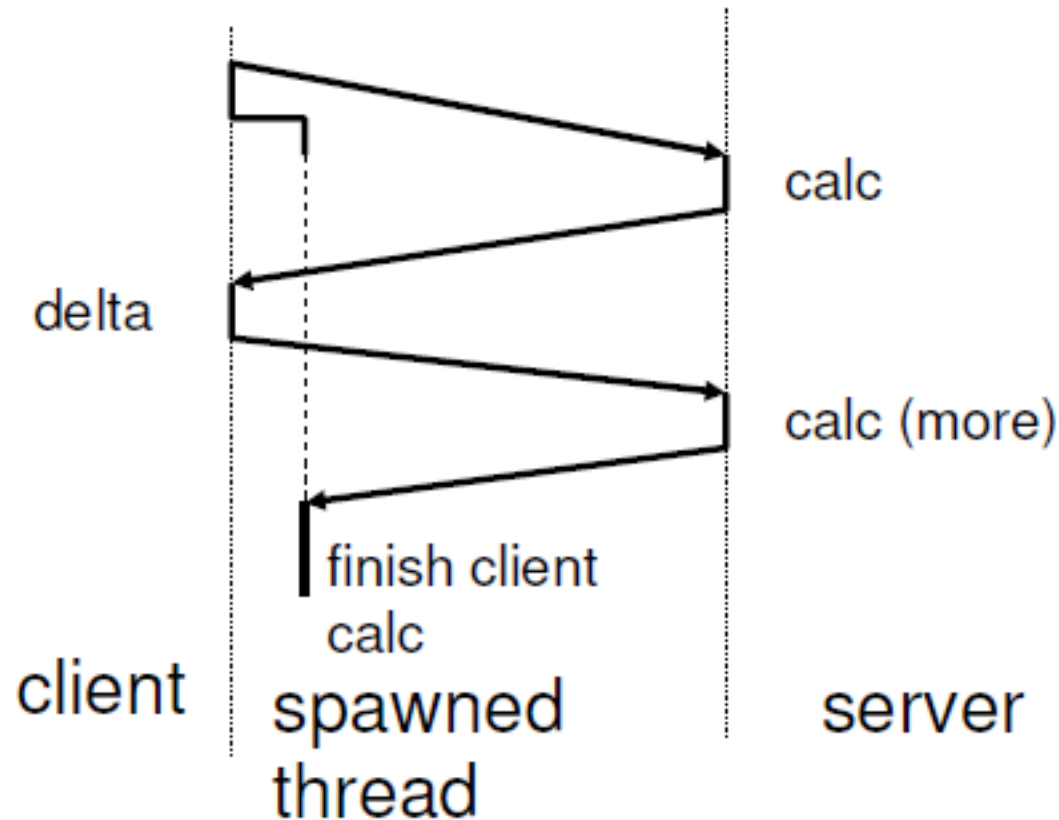
Solution – Use Thread

```
proc {ClientProc Msg}
  case Msg
  of work(?Z) then Y in
    {Send Server calc(10.0 Y Client)}
    thread Z=Y+100.0 end
  [] delta(?D) then D=1.0
  end
end
```



add thread to allow
client to proceed.

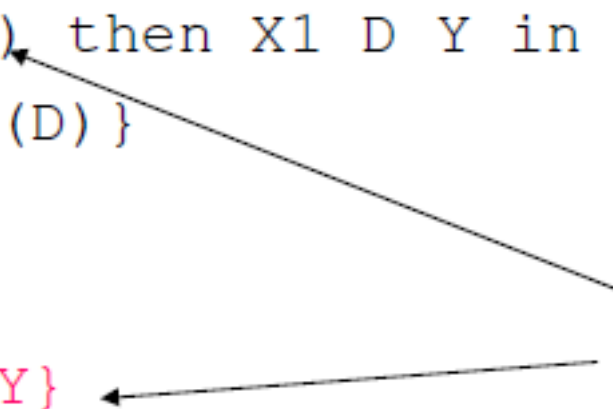
RMI with Callback (using thread)



RMI with Callback (using continuation)

- Possible to avoid thread.

```
proc {ServerProc Msg}
  case Msg
  of calc(X Client Cont) then X1 D Y in
    {Send Client delta(D)}
    X1=X+D
    Y = X * X + 1.0
    {Send Client Cont#Y}
  end
end
end
Server={NewAgent ServerProc}
```




continuation

Solution – Using Continuation Record

```
proc {ClientProc Msg}
  case Msg
  of work(?Z) then Y in
    {Send Server calc(10.0 Client cont(Z))}
    [] cont(Z)#Y then Z=Y+100.0
    [] delta(?D) then D=1.0
  end
end
end
Client={NewAgent ClientProc}
{Browse {Send Client work($)}}
```

RMI with Callback

(using continuation record)



Erlang

Erlang

- Developed by Ericsson for telecoms application.
 - Features : fine grain parallelism, extreme reliability, hot code updates.
 - Functional core – dynamically typed strict functional language.
 - Message-passing extension – processes communicate by sending messages asynchronously in FIFO order.
-

Functions in Erlang

- Uses pattern-matching and Prolog syntax

```
factorial(0) -> 1;
```

```
factorial(N) when N>0 -> N*factorial(N-1).
```


Pattern-Matching with Tuple

```
area({square, Side}) -> Side*Side;
```

```
area({rectangle, X, Y}) -> X*Y;
```

```
area({circle, R}) -> 3.14159*R*R;
```

```
area({triangle, A, B, C}) -> ... ;
```



tuple

Concurrency and Message Passing

- `spawn (M, F, A)` creates a new process and returns its `Pid`. Note that M-module, F-initial function, A-argument list.
 - Send operation (written as `Pid!msg`) is an asynchronous message sending.
 - `receive` operation removes message from a mailbox. It uses pattern-matching to select messages for removal
-

An Erlang Process

```
-module(areaserver)
```

```
-export([start/0, loop/0])
```

spawn

```
start() -> spawn(areaserver, loop, []).
```

```
loop() -> receive
```

receive

```
{From, Shape} ->
```

```
    From!area(Shape),
```

```
    loop()
```

```
end.
```

send

Receive Construct

```
receive
  Pattern1 [when Guard1] -> Body1;
  :
  PatternN [when GuardN] -> BodyN;
  [after Expr -> BodyT; ]
end
```

This expression blocks until a message matching one of patterns arrives or when timeout occurs

Summary

- Stream Object
 - Thread Module and Composition
 - Soft Real-Time Programming
 - Agents and Message Passing
 - Protocols
 - Erlang
-